

CSCI 270 Lecture 15: Optimal Caching

A cache can store n items. There is a sequence of m requests d_1, d_2, \dots, d_m known in advance. If an item is requested which is not in the cache, it must be brought into the cache, resulting in a **cache miss**. You may only bring something into the cache when it is requested. The goal is to minimize the number of cache misses.

$k = 2$, initial contents = 'ab'

requests: a b c b c a a b

- What greedy criteria might work for this problem?
- Can you find any counter-examples to these algorithms?
- How is this problem different than the standard version of Caching?
- What must happen at request $k + 1$?

<i>FIF</i>	...				
	Request k	f	o	e	...
	Request $k + 1$	d_{k+1}	o	e	...
	...				

<i>OPT</i>	...				
	Request k	f	o	e	...
	Request $k + 1$	f	d_{k+1}	e	...
	...				

At request $k + 1$, *FIF* kicks out f , while *OPT* kicks out o . Otherwise, the cache contents are completely identical.

Events that can cause our different cache contents to be relevant:

1. *OPT* kicks out f .
2. Cache miss on o .
3. Cache miss on f . This can't be the first event, because there must first be a cache miss on o .

We will refer to the first of these events as having occurred at request j .

Case 1: *OPT* removes f .

<i>OPT</i>	...				
	Request k	f	o	e	...
	Request $k + 1$	f	d_{k+1}	e	...
	...				
	Request $j - 1$	f	d_{k+1}	e	...
	Request j	d_j	d_{k+1}	e	...
	...				

We want to make an exchange, thereby transforming OPT into OPT' . We want OPT' to have the same number of cache misses as OPT . We also want OPT' to be identical to FIF through $k + 1$ requests.

OPT'	...				
	Request k	f	o	e	...
	Request $k + 1$	d_{k+1}	o	e	...
	...				
	Request $j - 1$	d_{k+1}	o	e	...
	Request j				...
	...				

What should OPT' do at request j ?

Case 2: There is a request on o

OPT	Request k	f	o	e	...
	Request $k + 1$	f	d_{k+1}	e	...
	...				
	Request $j - 1$	f	d_{k+1}	e	...
	Request j	f	d_{k+1}	o	...
	...				

OPT has to kick something out for o_2 , we will say it kicks out the here-to-fore unmentioned e .

We want to make an exchange, thereby transforming OPT into OPT' . We want OPT' to have the same number of cache misses as OPT . We also want OPT' to be identical to FIF through $k + 1$ requests.

OPT'	Request k	f	o	e	...
	Request $k + 1$	d_{k+1}	o	e	...
	...				
	Request $j - 1$	d_{k+1}	o	e	...
	Request j				...
	...				

- What would we like OPT' to do this at this request?
- Does this break any rules?

We will delay the exchange until an event occurs. OPT' has e , but OPT has f . The cache contents are otherwise identical. The possible events are:

- Something (x) is removed for e . We now plan to kick out x at the next event, instead of e . We avoid a cache miss, meaning that OPT wasn't even optimal.
- f is removed for something. We kick out e instead. We avoid a cache miss.
- f is requested. Now we kick out e as originally planned. We simply delayed our cache miss: it occurs for both solutions, but at different times.