A close-up photograph of a chessboard with several dark wooden chess pieces standing and one white piece lying on its side. The text "Divide & Conquer" is overlaid in white. A thin white line starts from the left, curves around the text, and extends towards the bottom right. In the bottom right corner, there are three white chevron symbols pointing to the right.

Divide & Conquer

Divide and Conquer

Divide and Conquer is a **runtime improvement technique**, that you can apply on top of existing algorithms of any type (Brute Force, Greedy, Dynamic Programming, etc.)

- MergeSort and QuickSort are the standard examples of Divide and Conquer.
- Take a large problem, and break it up into two or more smaller pieces (divide)
- Recursively solve the smaller pieces, by continuing to break them up until you reach constant-sized base cases.
- Merge your solved pieces together to get the overall solution (conquer)

Recursion

Dynamic Programming recursion is **completely different** than Divide & Conquer recursion.

- DP recursion is sequential (bite-sized decisions)
- D&C recursion is parallelized, with each subproblem independent of the others.
- DP recursion can be turned into a much more efficient iterative algorithm.
- D&C recursion is hard to unroll, and does not net an asymptotic improvement to runtime.



Counting Inversions

You are writing software for Pandora²⁷⁰.

- Each user ranks some number n of songs, from favorite to least favorite.
- You want to find users with similar tastes in music, so that you can recommend songs.
- You will take the songs that two users have in common, and compare their relative rankings.
- User 1 ranks 5 songs, conveniently named 1, 2, 3, 4, 5, ranked in that order.
- User 2 ranks them 1, 3, 4, 2, 5. How similar are these rankings?
- In one sense, they only agree on their favorite and least favorite songs, so they are 40% similar.
- Why is that an unfair metric?

Inversions

1 2 3 4 5

1 3 4 2 5

How many songs do we actually disagree on?

- Just song 2, which is off by two places.

An **inversion** is a pair of songs i and j , such that $i < j$, but song j ranked better than song i .

How many inversions could there be for n songs?

- $C(n,2) = \frac{n(n-1)}{2}$, or 10 when $n=5$.

What are our inversions for the above example?

- $\{2,3\}$ and $\{2,4\}$.

We have 2 out of 10 possible inversions, so we should say that our lists are **80%** similar.



Brute-Force

We want to count the number of inversions in a list of n numbers (or n songs).

- There are $\Theta(n^2)$ possible inversions, so we could count each inversion in quadratic time by checking each one manually.

Maybe we can do better, using Divide & Conquer!

Consider the following ranking:

1 5 4 8 10 2 6 9 12 11 3 7

Divide in $O(1)$ time:

1 5 4 8 10 2 | 6 9 12 11 3 7

Recursively solve each piece:

5 inversions | 8 inversions

Then merge our two solutions together.

A First Attempt

Our recursive call needs a base case.

- When you're down to a single song, return 0.

If the left subproblem returns 5 inversions, and the right subproblem returns 8 inversions, do we just return 5+8 for our total inversions?

- We also need to count every inversion spanning the division (which turns out to be 9 inversions).

What's the recurrence relation for our divide & conquer algorithm?

- $f(n) = 2f\left(\frac{n}{2}\right) + \Theta(n^2)$

$f(n) = \Theta(n^2)$

- Well that was useless...

You tried your best
and you failed miserably.



The lesson is,
never try.

CRITICAL LAYOUTS

A Second Attempt

Often your first attempt at a D&C algorithm will net no improvement. It **will** identify which part of the algorithm needs to be improved.

- The combine phase is the problem.

How would we need to store the songs in each list, so as to count the inversions more easily?

- How does MergeSort merge so easily?

We will both count inversions **and** sort the list.

- We are not losing any information, because we've already counted the inversions we're eliminating when sorting.

The Algorithm

1 5 4 8 10 2 6 9 12 11 3 7

1 5 4 8 10 2 | 6 9 12 11 3 7

5 inversions | 8 inversions

1 2 4 5 8 10 | 3 6 7 9 11 12

2 4 5 8 10 6 7 9 11

1 2 3 4 5 6 7 8 9 10 11 12

+4 +2+2 +1

When pulling from the right list, we check how many numbers are remaining in the left list ($n-i+1$, where n is the total number in the left list, and i is our iterator).

- In linear time, we both merged the two lists, and counted up the 9 inversions.

This is the MergeSort variant for counting inversions, and thus achieves $n \log n$: BubbleSort is the traditional n^2 algorithm for this.

All-Pairs Shortest Paths Strikes Back

$ASP[i,x,z]$ stores the length of the shortest path from x to z , using $\leq i$ edges.

- $ASP[i,x,x] = 0$
- $ASP[0,x,z] = \infty$
- $ASP[i,x,z] = \min_{\langle x,y \rangle \in E} (c_{\langle x,y \rangle} + ASP[i-1, y, z])$

For $i = 0$ to $n-1$

For all nodes z

For all nodes x

Calculate $ASP[i,x,z]$

- Answers at $ASP[n-1]$

Runtime = $\Theta(mn^2)$

- In Dynamic Programming, we figured out which node to visit next. What would we do for Divide & Conquer?



All-Pairs Shortest Paths

We will figure out which node to visit in the middle.

It's still, at its core, Dynamic Programming, as we will try all possible middle-nodes and take the best.

- $ASP[i,x,z] = \min_{y \in V} (ASP[\lfloor \frac{i}{2} \rfloor, x, y] + ASP[\lceil \frac{i}{2} \rceil, y, z])$
- $ASP[1,x,z] = c_{\langle x,z \rangle}$
- $ASP[i,x,x] = 0$

For $i = 1$ to $n-1$

For all nodes x

For all nodes z

Calculate $ASP[i,x,z]$

Runtime = $\Theta(n^4)$

...oops

A Second Attempt

What values of i do we actually need to calculate?

For $i = 0$ to $\log n$

For all nodes x

For all nodes z

Calculate $ASP[i, x, z]$


$ASP[i, x, z]$ stores the length of the shortest path from x to z , using no more than 2^i edges.

- Answers are at $ASP[2^{\lceil \log n \rceil}]$

Runtime = $\Theta(n^3 \log n)$

- This is usually better than running Bellman-Ford n times, but depends on how dense the graph is.

Integer Multiplication



Presumably you recall how to multiply numbers:

$$\begin{array}{r} 12 \\ \times 13 \\ \hline 36 \\ \underline{120} \\ 156 \end{array}$$

The same idea works for numbers in binary:

$$\begin{array}{r} 1100 \\ \times 1101 \\ \hline 1100 \\ 110000 \\ \underline{1100000} \\ 10011100 \end{array}$$

Elementary Math

What is the runtime to multiply two n -bit integers, using this method?

- $\Theta(n^2)$

We will try to improve this, using Divide & Conquer.

- We want to multiply two n -bit integers, X and Y .
- Let x_F be the first $\frac{n}{2}$ bits of X , and x_L be the last bits.
- Similarly for y_F and y_L .
- So, $x = x_F \cdot 2^{\frac{n}{2}} + x_L$ and $y = y_F \cdot 2^{\frac{n}{2}} + y_L$
- $x \cdot y = x_F \cdot y_F \cdot 2^n + (x_F \cdot y_L + y_F \cdot x_L) \cdot 2^{\frac{n}{2}} + x_L \cdot y_L$

Our base case is multiplying two 1-bit numbers.

What's the recurrence relation?

- $f(n) = 4 f(\frac{n}{2}) + \Theta(n)$



A Second Attempt

Where's the problem in our recurrence relation?

- The 4. We have too many subproblems.

How can we do this with only 3 subproblems?

- Hint: our first multiplication will be $(x_F + x_L) \cdot (y_F + y_L)$
- $M_2 = x_F \cdot y_F$
- $M_3 = x_L \cdot y_L$
- $x \cdot y = M_2 \cdot 2^n + (M_1 - M_2 - M_3) \cdot 2^{\frac{n}{2}} + M_3$

$$f(n) = 3 f\left(\frac{n}{2}\right) + \Theta(n)$$

- $f(n) = \Theta(n^{\log 3}) \approx n^{1.58}$

Sequence Alignment

The **edit distance** between two strings $X = x_1 \dots x_n$ and $Y = y_1 \dots y_m$ is the minimal number of changes to transform X into Y .

- $SA[i,j]$ is the edit distance of $X = x_1 \dots x_n$ and $Y = y_1 \dots y_m$.
- $SA[i, m+1] = n+1-i$
- $SA[n+1, j] = m+1-j$
- $SA[i, j] = SA[i+1, j+1]$, if $x_i = y_j$.
- $SA[i, j] = 1 + \min(SA[i, j+1], SA[i+1, j], SA[i+1, j+1])$, otherwise
- Runtime and space are both $\Theta(mn)$

Improving the space requirements

To calculate the matrix, we only need to remember the current and previous column.

- Therefore, we can reduce the memory requirements to $\Theta(m+n)$
- The cost is that we cannot reconstruct the solution.
- We will attempt to get the best of all worlds, using Divide & Conquer: reconstructing the answer in $\Theta(mn)$ time and $\Theta(m+n)$ space

The high level idea

Some prefix of Y will match with the first half of X. The rest of Y will match with the second half of X.

- We don't know where the optimal split point of Y is, so, dynamic programming-style, we will try all possible splits and take the best one.
- We will run Sequence Alignment, using our memory-saving technique, on the first half of X and all prefixes of Y.
- We will also run it on the second half of X and all suffixes of Y.
- Then we will inspect the results and figure out where the optimal split point is.

The high level idea: visually

Prefix of Y	First half of X	Suffix of Y	Second half of X	Total
ε	5	$y_1 \dots y_m$	4	9
y_1	4	$y_2 \dots y_m$	3	7
$y_1 y_2$	3	$y_3 \dots y_m$	2	5
$y_1 y_2 y_3$	2	$y_4 \dots y_m$	4	6
...
$y_1 \dots y_{m-1}$	3	y_m	5	8
$y_1 \dots y_m$	4	ε	6	10

Divide & Conquer

We will have determined the total edit distance, as well as the optimal split point, using only $\Theta(m+n)$ space.

$x_1x_2x_3x_4$	$x_5x_6x_7x_8$
$y_1y_2y_3$	$y_4y_5y_6y_7y_8y_9$

Now we repeat the process on the left-subproblem, and the right-subproblem

x_1x_2	x_3x_4	x_5x_6	x_7x_8
y_1	y_2y_3	y_4y_5	$y_6y_7y_8y_9$

We repeat until we are matching substrings of Y with individual characters of X , at which point we have the full reconstruction.

x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8
y_1	ϵ	y_2	y_3	y_4y_5	ϵ	y_6y_7	y_8y_9

Sequence Alignment on the last half of X, and suffixes of Y

Run Sequence Alignment on the last half of X and **all** of Y (using our space-saving technique).

Runtime
= $\Theta(mn)$

	$i = \frac{n}{2}$	$i = \frac{n}{2} + 1$...	$i = n$
$j=1$	4	5	We	are
$j=2$	3	4	not	saving
...	these	columns,
$j=m$	5	4	to	save
$j=m+1$	6	5	space	!

The first column stores the edit distances of all possible suffixes of Y with the last half of X.

Sequence Alignment on the first half of X, and prefixes of Y

Run Sequence Alignment on the first half of X and **all** of Y (using our space-saving technique).

	$i = 1$	$i = 2$...	$i = \frac{n}{2} - 1$
$j=1$	4	5	We	are
$j=2$	3	4	not	saving
...	these	columns,
$j=m$	5	4	to	save
$j=m+1$	6	5	space	!

Does the first column store the edit distances of all possible prefixes of Y with the first half of X?

- Nope, it's still suffixes of Y.

Sequence Alignment on the first half of X, and prefixes of Y, Take two

Run Sequence Alignment on the first half of X and the reverse of Y (using our space-saving technique).

	$i = 1$	$i = 2$...	$i = \frac{n}{2} - 1$
$j = m$	4	5	We	are
$j = m - 1$	3	4	not	saving
...	these	columns,
$j = 1$	5	4	to	save
$j = 0$	6	5	space	!

Does this work now?

- Nope, now we're matching the first character of X with the last character of Y.

Sequence Alignment on the first half of X, and prefixes of Y, Take three

Run Sequence Alignment on the first half of X (reversed) and the reverse of Y (using our space-saving technique).

	$i = \frac{n}{2} - 1$	$i = \frac{n}{2} - 2$...	$i = 1$
$j = m$	4	5	We	are
$j = m - 1$	3	4	not	saving
...	these	columns,
$j = 1$	5	4	to	save
$j = 0$	6	5	space	!

Now this works!

Analysis

Each individual call takes only $\Theta(m+n)$ space.

- Additionally, we need to store the reconstruction, but that also takes only $\Theta(m+n)$ space.

Our recurrence relation is

$$f(n, m) = f\left(\frac{n}{2}, k\right) + f\left(\frac{n}{2}, m-k\right) + \Theta(mn).$$

- The next level of recursion will get runtime

$$\frac{n}{2}k + \frac{n}{2}(m - k) = \frac{n}{2}m$$


- Each successive level divides the runtime in half, so the total runtime is $\Theta(mn)$

Closest Points on a Plane

Given n points on a plane (specified by their x and y coordinates), find the pair of points with the smallest Euclidean distance between them.

What would be the runtime for a Brute-Force algorithm?

- $\Theta(n^2)$



Divide & Conquer

How should we divide the problem?

- We can divide at the median x-coordinate, so that half of the points are on the left, and half of the points are on the right.
- If we also try to split on the y-coordinate, there is no guarantee we will be able to split the problem into equal quarters.



Recursion

We will continually split the problem until we have 2-3 points, and then we will brute-force identify the closest pair of points among them.


- So, if I've found the closest pair of points on the left, with distance δ , and the closest pair of points on the right, with distance ε , do I just return $\min(\delta, \varepsilon)$?
- We also need to check points that span the divide. If we brute force this, we'll get the following recurrence:
- $f(n) = 2f(\frac{n}{2}) + \Theta(n^2)$

The Combine Phase, Take Two

Do I really need to compare all pairs of points that span the divide?

- Assume δ is the smallest distance over both subproblems. Then we only need to look at points that are within δ of the dividing line.


Does this improve our runtime?



The Combine Phase, Take Three

Do I really need to compare all pairs of points within δ of the dividing line?

- For each point, I only need to look at the points on the other side that are within δ of the given y-coordinate.
- Does this improve the runtime?
- It does! Since the smallest pair of points on the right side are at least δ apart, we can only fit a constant number of points in the rectangle!
- The new recurrence is $f(n) = 2f(\frac{n}{2}) + \Theta(n)$



Details of Implementation

1. Sort the list of points (list 1) by x-coordinate:
 $\Theta(n \log n)$
2. Sort the list of points (list 2) by y-coordinate:
 $\Theta(n \log n)$
3. Use list 1 to figure out the division in constant time, and recursively call the algorithm.
4. Walk through list 2, and pull out the points that are within δ of the dividing line (list 3).
5. For each point in list 3, compare it with each successive point in list 3 until you are over δ away from that point: this is a constant number of comparisons.



Matrix Multiplication

The standard way to multiply matrices:

$$\begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix} \begin{pmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{pmatrix} = \begin{pmatrix} a_{11}b_{11} + a_{12}b_{21} & a_{11}b_{12} + a_{12}b_{22} \\ a_{21}b_{11} + a_{22}b_{21} & a_{21}b_{12} + a_{22}b_{22} \end{pmatrix}$$

$$c_{ij} = \sum_k a_{ik} b_{kj}$$

Using this method, how long will it take to multiply two n by n matrices?

- $\Theta(n^3)$

Divide & Conquer

We can divide our n by n matrices into 4 quadrants:

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix} \begin{pmatrix} E & F \\ G & H \end{pmatrix} = \begin{pmatrix} AE + BG & AF + BH \\ CE + DG & CF + DH \end{pmatrix}$$

Then we need to figure out how to multiply two $\frac{n}{2}$ by $\frac{n}{2}$ matrices together (AE, BG, AF, BH, CE, DG, CF, DH).

We continue to break these down until we're multiplying scalars.

What's the recurrence relation?

- $f(n) = 8f\left(\frac{n}{2}\right) + \Theta(n^2)$

Divide & Conquer, Take Two

It is possible to reduce the number of multiplies:

1. $M_1 = A(F-H)$

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix} \begin{pmatrix} E & F \\ G & H \end{pmatrix} =$$

2. $M_2 = (A+B)H$

$$\begin{pmatrix} AE + BG & AF + BH \\ CE + DG & CF + DH \end{pmatrix} =$$

3. $M_3 = (C+D)E$

4. $M_4 = D(G-E)$

$$\begin{pmatrix} M_5 + M_4 - M_2 + M_6 & M_1 + M_2 \\ M_3 + M_4 & M_1 + M_5 - M_3 - M_7 \end{pmatrix}$$

5. $M_5 = (A+D)(E+H)$

$$f(n) = 7f\left(\frac{n}{2}\right) + \Theta(n^2)$$

6. $M_6 = (B-D)(G+H)$

$$f(n) = \Theta(n^{\log 7}) \approx \Theta(n^{2.81})$$

7. $M_7 = (A-C)(E+F)$

For $n = 2500$, this has a speedup factor of 8

Improvements

This is known as Strassen Multiplication, and was discovered in 1969.

- In 1971, Hopcroft and Kerr showed that 6 multiplies is impossible
- In 1980, Pan showed that if you split the matrix up into 70^2 sectors, you can achieve 143640 multiplies $\approx \Theta(n^{2.80})$
- In December 1979, we got a general $O(n^{2.521813})$ algorithm.
- In January 1980 this was improved to $O(n^{2.521801})$
- In 1987, Coppersmith and Winograd achieved $O(n^{2.376})$
- There have been small improvements starting in 2010
- Hypothesis: $O(n^{2+\epsilon})$ is possible

Take-Home Practice



Chapter 5



Exercises 1, 2, 3, 6, 7